
CS 7180 Milestone 2

Nalin Gupta
NUID: 001497315

Christopher Botica
NUID: 001726854

Tyler Brown
NUID: 001684955

1 Introduction

Images taken in low-light conditions are often too dark, noisy, and distorted to be used in industrial purposes. We propose a deep-learning model that processes low-light images to improve image brightness and increase their overall quality. The problem with imaging in low-light conditions is challenging due to low-photon count and low Signal-to-Noise (SNR) ratio. These yield very dark and noisy images. The most common technique to overcome this problem is long exposure shot. However, this method yields blurry images with the slightest camera shake or object motion[1]. Common post-processing techniques brighten the image at the expense of image quality. Being able to “see in the dark” provides a number of real-world benefits such as photography, computer vision, and social networking.

2 Related Work

In the past, the problem of enhancing low light images has been tackled via noise reduction. This noise becomes dominant especially in low-light images due to low SNR. Remez et. al. proposed a deep CNN for noise reduction under the assumption that this low-light noise belongs to a Poisson distribution [2]. They used images from ImageNet [3] as their ground truth data and added synthetic Poisson noise to simulate corrupted images. Even though their model outperform the state-of-the art de-noiser “BM3D”, it does not scale well to real world images, due to their underlying assumptions. Furthermore, their model only denoises images but does not brighten them. Motivated by these downfalls, Chen et. al., proposed an end-to-end CNN, “See-in-the-Dark” (SID), which brightens extremely low light images and removes noise without making any underlying assumptions [1]. However these advances come with the added expense of collecting large amounts of low light and bright light images. In the absence of a true vs noisy image dataset, the team captured scenes using various exposure times to generate true (bright light) and corrupted (low light) image pairs called “See-in-the-Dark Dataset” (SID Dataset¹). Furthermore, their model is camera specific and not easily generalizable.

We propose a transferable CNN for image brightening and denoising. Instead of training our model on actual true (bright light) and corrupted (low light) image pairs, we use images from the publicly available RAISE raw images dataset [4]. We use the RAISE raw images dataset as our true images and corrupt these by simulating low-light conditions for use in our model. We train our CNN on this synthetic data to obtain our initial model parameters. A small fraction of the real images paired from the SID dataset is then used in our transfer learning [5] approach to update our model parameters. These model parameters are updated for the particular camera used in the SID dataset. We then use this model to test on the SID Dataset. In addition, we aim to test various transfer learning approaches, such as the traditional transfer learning and zero shot learning [6, 7, 8].

The novelty of our approach stems from the idea of “more for less”. Our model drastically reduces the overhead costs of data collection by synthesizing readily available training data (RAISE raw images

¹<https://github.com/cchen156/Learning-to-See-in-the-Dark>

dataset). This is particularly beneficial in domains where collecting images pairs is expensive/time consuming.

3 Methods

Our contributions were to simulate the data and replicate their model using simulated data.

3.1 Dataset

3.2 Model

As illustrated in Figure 1, the traditional pipeline takes a corrupted image, and applies the following sequence of modules: Reduce Black Level, Denoising, White Balance, and Gamma Correction. The Black Level refers to the level of brightness at the darkest parts of the image, and is reduced by subtracting the minimum pixel value. Denoising is reduced using common algorithms such as BM3D. White Balance refers to the color balance in the image (i.e., white should be true white) and is corrected by re-balancing the intensities of each color RGB. Finally, Gamma Correction controls the overall brightness of the image. We synthetically generate corrupted images by applying the reverse of this pipeline. Gamma Distortion: decrease the brightness of the image, White Imbalance: skew the color-space by multiplying each level of RGB by a random weight, Poisson Noise: add Poisson noise to the image, Black Level: add a negative bias to the pixel values (i.e., random black level).

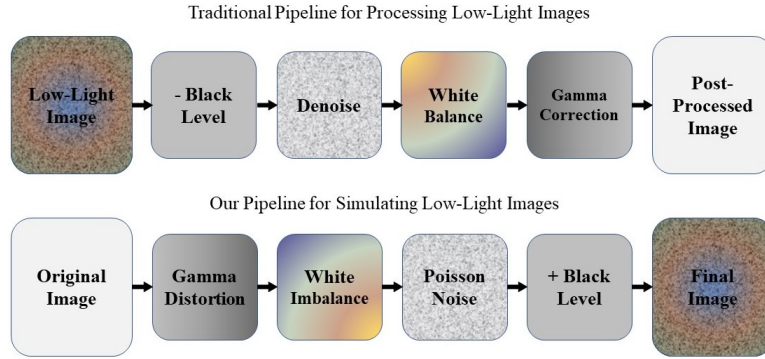


Figure 1: Top: Traditional Pipeline for processing low-light images. Bottom: Our pipeline to simulate low-light images based on the traditional, only in reverse.

Our model is based on the one developed for SID, with the addition of 2 fully-connected layers. Note that we may increase this if training runtime permits. Using the transfer learning approach, we will first train our model on the MIT dataset, then erase the learned weights from the last two fully-connected layers, and retrain on the SID dataset. This is highlighted in our model framework in Figure 2.

We tried to replicate their model to using other readily available data such as CIFAR10 [9] and ImageNet [3]. The CIFAR10 and ImageNet data is augmented to simulate properties of the images used by Chen et. al. (2018) [1]. The CIFAR10 and ImageNet datasets are distinctly different because they are not *raw* images. This means the dimensionality of CIFAR10 is $32 \times 32 \times 3$ rather than a raw image which could be $512 \times 512 \times 4$. We not only need to simulate images but also account for the *change in dimensionality* across image types.

3.3 Computational Resources

Using AWSEducate did not work for us. We were unable to create roles with IAM authentication so it's really hard or impossible to move data from a S3 Bucket to an EC2 Instance. We tried to create a *p3.8xlarge* instance but these instances are not allowed even though they are listed. Using regular AWS does work but is costly. Ran a single AWS EC2 *p3.8xlarge* instance with 32 CPU, 244 GB of Memory, 4 Tesla V100 GPUs, and 64 GPU Memory. This costs \$12.24 an hour. This is the amount

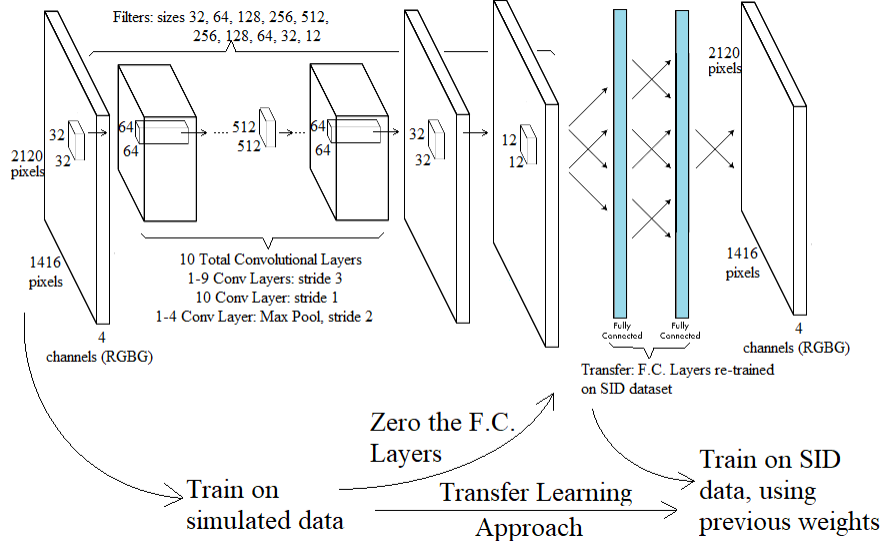


Figure 2: Our proposed model framework

of GPU Memory requested by the paper authors as a minimum amount.

We have also been able to use Google Cloud for some of our workload. This resource is well integrated with Tensorflow. Leading up to Milestone 2, we extensively used Google Cloud Compute with 2 CPU and 7.5GB of memory for debugging and working with image datasets; primarily CIFAR10 and Tiny ImageNet.

In regards to training our model, one of our group members has also received permission to use computational resources from Lincoln Laboratory's Super Computer (LLSC) for our project. We are currently waiting for complete access and expect to be able to use it in the coming days.

3.4 Low Light Simulation

TODO: some theory, maybe Chris?



Figure 3: SID Images in Normal and Low Light Conditions

4 Experiment

We experimented on our contributions

4.1 Model Implementation

We tried replicating their model and writing a model in parallel to replicate their functionality. We also tried using the CIFAR10 and ImageNet datasets as input. The primary issue when replicating their model related to dimensionality. The RGB images in CIFAR10 and ImageNet have $m \times n \times 3$ dimensions. Raw images have $m \times n \times 4$ dimensions.

We were able to confirm that output dimensions from Chen et. al. [1] follow the pattern $2m \times 2n \times 3$ irrespective of the image size or type used. During our experiment, we focused our debugging efforts on understanding whether the dimensionality of a raw image as an input generated incompatible dimensions with the target vector after model output. This error resulted in our replicated model not being able to compute the cost function.

In terms of techniques used, their model is mostly written in a higher level Tensorflow contribution module called “slim”. The “slim” contribution module does not appear to have any documentation and requires one to read the module code embedded within Tensorflow. We initially tried writing a model in parallel using Keras but found that moving from Keras objects back down to lower level Tensorflow for some blocks within forward propagation to not be documented well. We then rewrote the Chen et. al. [1] model using only low level Tensorflow code. This approach has the disadvantage of requiring the dimensions of all weight matrices to be specified and initialized in a separate function we called ‘initialize_parameters()’.

These efforts allowed us to understand the dimensionality of both models at each of the 10 blocks within their forward propagation. We also experimented with pooling in regards to window size. We also experimented with strides for convolutional layers with each block. We also added up to three additional blocks with different parameters in an attempt to get the output dimensions to match the dimensionality of our input image.

At this point, we have not been able to use RGB images, having $m \times n \times 3$ dimensions, to work with Chen et. al. [1] model. It is important for us to develop this capability in order to qualitatively verify that we are correctly augmenting image data. We are unable to generate images in raw format because it would require us to know proprietary information about how a camera generated an image.

In order to generate output for this milestone, we modified images from Chen et. al. [1], such that

$$\begin{aligned} X &= g(y_{train}) \\ Y &= y_{train} \end{aligned}$$

where function g is an application of artificial noise to a clean image in y_{train} which is then trained against a corresponding image where g has not been applied.

4.2 State of the Art Model Results

The model used by Chen et. al. [1] trains on raw images for 4,000 epochs. It beats the current state of the art denoising technique, BM3D. An example of a result for this model is shown in Figure 4. A drawback of this model is that it learns a separate model for each device used to capture the raw images.



Figure 4: State of the Art Model Results

We can see in Figure 4, a nearly dark image in sub-figure (a) is converted to an image shot in normal light conditions; shown in sub-figure (b).

4.3 Our Model Results

Currently, our model has been trained on a subset of simulated lowlight images using the SID dataset as mentioned in Section 4.1. We ran the model for 25 epochs due to limited resources. However, we were still able to obtain images which showed some enhancement. We are able to show some initial results but they are not of the same quality compared to state-of-the-art. [1].




Figure 5: Predicted Results

Figure 5 shows the steps followed by our approach. Sub-figure (a) is the original image, sub-figure (b) is the simulated low light image described in Section 3.4, and sub-figure (c) shows the result of our model. We can clearly observe from Figure 5 that our model was able to identify the colors while simultaneously brightening the image.




Figure 6: Loss vs. Epochs

Figure 6 plots our loss function at each epoch. We can see a downward trend in the loss function output. However, we can also see that our model is having trouble converging after 25 epochs. We believe this problem can be solved when we use the entire dataset to train the model.

5 Discussion

6 Appendix

References

- [1] Chen Chen, Qifeng Chen, Jia Xu, and Vladlen Koltun. Learning to see in the dark. *arXiv preprint arXiv:1805.01934*, 2018.
- [2] Tal Remez, Or Litany, Raja Giryes, and Alex M Bronstein. Deep convolutional denoising of low-light images. *arXiv preprint arXiv:1701.01687*, 2017.

- [3] J. Deng, W. Dong, R. Socher, L.-J. Li, K. Li, and L. Fei-Fei. ImageNet: A Large-Scale Hierarchical Image Database. In *CVPR09*, 2009.
- [4] Duc-Tien Dang-Nguyen, Cecilia Pasquini, Valentina Conotter, and Giulia Boato. Raise: A raw images dataset for digital image forensics. In *Proceedings of the 6th ACM Multimedia Systems Conference*, MMSys '15, pages 219–224, New York, NY, USA, 2015. ACM.
- [5] Ian Goodfellow, Yoshua Bengio, and Aaron Courville. *Deep Learning*. MIT Press, 2016. <http://www.deeplearningbook.org>.
- [6] Hugo Larochelle and Y Bengio. Classification using discriminative restricted boltzmann machines. pages 536–543, 01 2008.
- [7] Mark Palatucci, Dean Pomerleau, Geoffrey E Hinton, and Tom M Mitchell. Zero-shot learning with semantic output codes. In Y. Bengio, D. Schuurmans, J. D. Lafferty, C. K. I. Williams, and A. Culotta, editors, *Advances in Neural Information Processing Systems 22*, pages 1410–1418. Curran Associates, Inc., 2009.
- [8] Richard Socher, Milind Ganjoo, Hamsa Sridhar, Osbert Bastani, Christopher D. Manning, and Andrew Y. Ng. Zero-shot learning through cross-modal transfer, 2013.
- [9] Alex Krizhevsky, Vinod Nair, and Geoffrey Hinton. Cifar-10 (canadian institute for advanced research).